

GPU Programming

Rupesh Nasre.

<http://www.cse.iitm.ac.in/~rupesh>

IIT Madras
July 2017

Let's Sync-up!

- Print numbers 1 to 10 in sequence.
- Launch kernel with 3 threads.
- Each kernel prints threadIdx.x modulo 3.

```
__global__ void onetoten() {  
    for (int ii = 0; ii < 10; ++ii)  
        if (ii % 3 == threadIdx.x)  
            printf("%d: %d\n", threadIdx.x, ii);  
}
```

```
__global__ void onetoten() {  
    __shared__ int n;  
    n = 0;  
    __syncthreads();  
    while (n < 10) {  
        if (n % 3 == threadIdx.x) {  
            printf("%d: %d\n", threadIdx.x, n);  
            ++n;  
        }  
        __syncthreads();  
    }  
}
```

Let's Sync-up!

- Print numbers 1 to 10 in sequence.
- Launch kernel with 3 threads.
- Each kernel prints threadIdx.x modulo 3.

```
__device__ volatile int n;

__global__ void onetoten() {
    n = 0;
    __syncthreads();
    while (n < 10) {
        if (n % 3 == threadIdx.x) {
            printf("%d: %d\n", threadIdx.x, n);
            ++n;
        }
    }
}
```

```
__global__ void onetoten() {
    volatile __shared__ int n;
    n = 0;
    __syncthreads();
    while (n < 10) {
        if (n % 3 == threadIdx.x) {
            printf("%d: %d\n", threadIdx.x, n);
            ++n;
        }
    }
}
```

Let's Sync-up!


- Print numbers 1 to 10 in sequence.
- Launch kernel with 3 threads.
- Each kernel prints `threadIdx.x` modulo 3.

```
__global__ void onetoten() {  
    __shared__ unsigned int n;  
    n = 0;  
    __syncthreads();  
  
    while (n < 10) {  
        int oldn = atomicInc(&n, 100);  
        if (oldn % 3 == threadIdx.x) {  
            printf("%d: %d\n", threadIdx.x, oldn);  
        }  
    }  
}
```

Note that some of these codes are faulty.

Let's Compute the Shortest Paths

- You are given an input graph of India, and you want to compute the shortest path from Nagpur to every other city.
- Assume that you are given a GPU graph library and the associated routines.



```
__global__ void dsssp(Graph g, unsigned *dist) {  
    unsigned id = ...  
    for each n in g.allneighbors(id) {    // pseudo-code.  
        unsigned altdist = dist[id] + weight(id, n);  
        if (altdist < dist[n]) {  
            dist[n] = altdist;  
        }  
    }  
}
```

What is the error in this code?

Data Race

- A datarace occurs if *all* of the following hold:
 1. Multiple threads
 2. Common memory location
 3. At least one write
 4. Concurrent execution
- Ways to remove datarace:
 1. Execute sequentially
 2. Privatization / Data replication
 3. Separating reads and writes by a barrier
 4. Mutual exclusion

Classwork

- Is there a data race in this code?
- What does the code ensure?
- Can you ensure the same with barriers?
- Can you ensure the same with atomics?
- Generalize it for N threads.

T1	T2
<pre>flag = 1; while (flag) ; S1;</pre>	<pre>while (!flag) ; S2; flag = 0;</pre>

Synchronization

- Atomics
- Barriers
- Control + data flow
- ...

atomics

- Atomics are primitive operations whose effects are visible either none or fully (never partially).
- Need hardware support.
- Several variants: atomicCAS, atomicMin, atomicAdd, ...
- Work with both global and shared memory.

atomics

```
__global__ void dkernel(int *x) {  
    ++x[0];  
}  
...  
dkernel<<<1, 2>>>(x);
```

After dkernel completes,
what is the value of x[0]?

++x[0] is equivalent to:

Load x[0], R1

Increment R1

Store R1, x[0]

Time
↓

Load x[0], R1

Increment R1

Store R1, x[0]

Load x[0], R2

Increment R2

Store R2, x[0]

Final value stored in x[0] could be 1 (rather than 2).

What if x[0] is split into multiple instructions? What if there are more threads?


atomics

```
__global__ void dkernel(int *x) {  
    ++x[0];  
}  
...  
dkernel<<<1, 2>>>(x);
```

- Ensure all-or-none behavior.
 - e.g., `atomicInc(&x[0], ...)`;
- **dkernel**<<<K1, K2>>> would ensure `x[0]` to be incremented by exactly $K1 * K2$ – irrespective of the thread execution order.
 - When would this effect be visible?

Let's Compute the Shortest Paths

- You are given an input graph of India, and you want to compute the shortest path from Nagpur to every other city.
- Assume that you are given a GPU graph library and the associated routines.



```
__global__ void dsssp(Graph g, unsigned *dist) {
    unsigned id = ...
    for each n in g.allneighbors(id) {    // pseudo-code.
        unsigned altdist = dist[id] + weight(id, n);
        if (altdist < dist[n]) {
            dist[n] = altdist;    atomicMin(&dist[n], altdist);
        }
    }
}
```

Classwork

1. Compute sum of all elements of an array.
2. Find the maximum element in an array.
3. Each thread adds elements to a worklist.
 - e.g., next set of nodes to be processed in SSSP.

AtomicCAS

- Syntax: `oldval = atomicCAS(&var, x, y);`
- Typical usecases:
 - *Locks* (critical section processing)
 - *Single*
 - Other atomic variants


Classwork: Implement *single* with *atomicCAS*.

Barriers

- A barrier is a program point where all threads need to reach before any thread can proceed.
- End of kernel is an implicit barrier for all GPU threads (**global barrier**).
- There is no explicit global barrier supported in CUDA.
- Threads in a thread-block can synchronize using **__syncthreads()**.
- How about barrier within warp-threads?

Barriers

```
__global__ void dkernel(unsigned *vector, unsigned vectorsize) {  
    unsigned id = blockIdx.x * blockDim.x + threadIdx.x;  
    vector[id] = id; S1  
    __syncthreads();  
    if (id < vectorsize - 1 && vector[id + 1] != id + 1) S2  
        printf("syncthreads does not work.\n");  
}
```



Barriers

- `__syncthreads()` is not only about control synchronization, it also has data synchronization mechanism.
- It performs a memory fence operation.
 - A memory fence ensures that the writes from a thread are made visible to other threads.
 - `__syncthreads()` executes a fence for all the block-threads.
- There is a separate `__threadfence_block()` instruction also. Then, there is `__threadfence()`.
- *[In general]* A fence does not ensure that other thread will read the updated value.
 - This can happen due to caching.
 - The other thread needs to use `volatile` data.
- *[In CUDA]* a fence applies to both read and write.

Classwork

- Write a CUDA kernel to find maximum over a set of elements, and then let thread 0 print the value in the same kernel.
- Each thread is given `work[id]` amount of work. Find average work per thread and if a thread's work is above $\text{average} + K$, push extra work to a worklist.
 - This is useful for load-balancing.
 - Also called work-donation.

Synchronization

- Atomics
- Barriers
- **Control + data flow**
- ...


Initially, flag == false.

```
while (!flag) ;  
S1;
```

```
S2;  
flag = true;
```

Reductions

- What are reductions?
- Computation properties required.
- Complexity measures



Classwork: Write the reduction code.

Reductions

```
for (int off = n/2; off; off /= 2) {  
    if (threadIdx.x < off) {  
        a[threadIdx.x] += a[threadIdx.x + off];  
    }  
    __syncthreads();  
}
```

log(n) steps	Input:	4	3	9	3	5	7	3	2	n numbers
		7		12			12		5	<i>barrier</i>
				19					17	
	Output:					36				

Reductions

```
for (int off = n/2; off; off /= 2) {  
    if (threadIdx.x < off) {  
        a[threadIdx.x] += a[threadIdx.x + off];  
    }  
    __syncthreads();  
}
```

- Write the reduction such that thread i sums $a[i]$ and $a[i + n/2]$.
- Assuming each $a[i]$ is a character, find a concatenated string using reduction.
- String concatenation cannot be done using $a[i]$ and $a[i + n/2]$, but computing sum was possible; why?
- What other operations can be cast as reductions?

Prefix Sum

- Imagine threads wanting to push work-items to a central worklist.
- Each thread pushes different number of work-items.
- This can be computed using atomics or prefix sum (also called as *scan*).

Input: 4 3 9 3 5 7 3 2

Output: 4 7 16 19 24 31 34 36

OR

Output: 0 4 7 16 19 24 31 34

Classwork: Write the prefix-sum code.

Prefix Sum

```
for (int off = n/2; off; off /= 2) {  
    if (threadIdx.x < off) {  
        a[threadIdx.x] += a[threadIdx.x + off];  
    }  
    __syncthreads();  
}
```

```
for (int off = n; off; off /= 2) {  
    if (threadIdx.x < off) {  
        a[threadIdx.x] += a[threadIdx.x + off];  
    }  
    __syncthreads();  
}
```

Input: 4 3 9 3 5 7 3 2

Output: 4 7 16 19 24 31 33 35

OR

Output: 0 4 7 16 19 24 31 33

Prefix Sum

```
for (int off = n/2; off; off /= 2) {  
    if (threadIdx.x < off) {  
        a[threadIdx.x] += a[threadIdx.x + (n - off)];  
    }  
    __syncthreads();  
}
```

```
for (int off = 0; off < n; off *= 2) {  
    if (threadIdx.x > off) {  
        a[threadIdx.x] += a[threadIdx.x - off];  
    }  
    __syncthreads();  
}
```

Input: 4 3 9 3 5 7 3 2

Output: 4 7 16 19 24 31 33 35

OR

Output: 0 4 7 16 19 24 31 33

Prefix Sum

x_0	x_1	x_2	x_3	x_4	x_5	x_6	x_7
-------	-------	-------	-------	-------	-------	-------	-------

$\Sigma(x_0 \dots x_0)$	$\Sigma(x_0 \dots x_1)$	$\Sigma(x_0 \dots x_2)$	$\Sigma(x_0 \dots x_3)$	$\Sigma(x_0 \dots x_4)$	$\Sigma(x_0 \dots x_5)$	$\Sigma(x_0 \dots x_6)$	$\Sigma(x_0 \dots x_7)$
-------------------------	-------------------------	-------------------------	-------------------------	-------------------------	-------------------------	-------------------------	-------------------------


Input: 4 3 9 3 5 7 3 2

Output: 4 7 16 19 24 31 33 35

OR

Output: 0 4 7 16 19 24 31 33

Prefix Sum




Input: 4 3 9 3 5 7 3 2

Output: 4 7 16 19 24 31 33 35

OR

Output: 0 4 7 16 19 24 31 33

Prefix Sum



Input: 4 3 9 3 5 7 3 2

Output: 4 7 16 19 24 31 33 35

OR

Output: 0 4 7 16 19 24 31 33

Prefix Sum

```
for (int off = 1; off < n; off *= 2) {  
    if (threadIdx.x > off) {  
        a[threadIdx.x] += a[threadIdx.x - off];  
    }  
    __syncthreads();  
}
```

Datarace

```
for (int off = 0; off < n; off *= 2) {  
    if (threadIdx.x > off) {  
        tmp = a[threadIdx.x - off];  
        __syncthreads();  
        a[threadIdx.x] += tmp;  
    }  
    __syncthreads();  
}
```

Separating
R and W
in time

Prefix Sum

```
for (int off = 1; off < n; off *= 2) {  
    if (threadIdx.x >= off) {  
        tmp = a[threadIdx.x - off];  
    }  
    __syncthreads();  
  
    if (threadIdx.x >= off) {  
        a[threadIdx.x] += tmp;  
    }  
    __syncthreads();  
}
```


Homework: Can one of the barriers be avoided?

Application of Prefix Sum

- Assuming that you have the prefix sum kernel, insert elements into the worklist.
 - Each thread inserts $\text{nelem}[\text{tid}]$ many elements.
 - The order of elements is not important.
 - You are forbidden to use atomics.
- Computing cumulative sum
 - Histogramming
 - Area under the curve

Concurrent Data Structures

- Array
 - atomics for index update
 - prefix sum for coarse insertion
- Singly linked list
 - insertion
 - deletion [marking, actual removal]



Concurrent Data Structures


Type definition

```
struct node {  
    char item;  
    struct node *next;  
};
```

G->next = P2;
P1->next = G;

Sequential insert

How to execute the
two instructions
atomically?




Concurrent Linked List

Solution 1: Keep a lock with the list.

- Coarse-grained synchronization
- Low concurrency / sequential access
- Easy to implement
- Easy to argue about correctness

Classwork: Implement lock() and unlock().



lock() and unlock()

```
void lock(List &list) {  
    while (list.sema == 1)  
        ;  
    list.sema = 1;  
}  
void unlock(List &list) {  
    list.sema = 0;  
}
```

time
↓

T1	T2	T3
sema = 1		
-- CS --		
sema = 0		
	sema == 1	sema == 1
	sema = 1	sema = 1
	-- CS --	-- CS --

What is the problem here? How to fix it?

lock() and unlock()

```
void lock(List &list) {  
    atomicCAS(&list.sema, 0, 1);  
}  
void unlock(List &list) {  
    atomicCAS(&list.sema, 1, 0);  
}
```

```
void lock(List &list) {  
    do {  
        old = atomicCAS(&list.sema, 0, 1);  
    } while (old == 1);  
}  
void unlock(List &list) {  
    list.sema = 0;  
}
```


Concurrent Linked List

Solution 2: Keep a lock with each node.

- Fine-grained synchronization
- Better concurrency
- Moderately difficult to implement, need to finalize the supported operations
- Difficult to argue about correctness when multiple nodes are involved

Classwork: Implement insert().

Classwork: Check if two concurrent inserts work.




Concurrent Linked List

```
void insert(Node &prev, Node &naya) {  
    naya.next = prev.next;  
    prev.lock();  
    prev.next = naya;  
    prev.unlock();  
}
```

Classwork: Implement insert().

Classwork: Check if two concurrent inserts work.




Concurrent Linked List

```
void insert(Node &prev, Node &naya) {  
    naya.next = prev.next;  
    prev.lock();  
    prev.next = naya;  
    prev.unlock();  
}
```

Classwork: Implement insert().

Classwork: Check if two concurrent inserts work.




Concurrent Linked List

```
void insert(Node &prev, Node &naya) {  
    prev.lock();  
    naya.next = prev.next;  
    prev.next = naya;  
    prev.unlock();  
}
```

Classwork: Implement insert().

Classwork: Check if two concurrent inserts work.



Classwork: Now allow remove().


```
void insert(Node &prev, Node &naya) {
    prev.lock();
    naya.next = prev.next;
    prev.next = naya;
    prev.unlock();
}
```


```
void remove(Node &prev, Node &tbr) {
    prev.lock();
    tbr.lock();
    prev.next = tbr.next;
    // process tbr.
    tbr.unlock();
    prev.unlock();
}
```

If the order of locks is reversed?

Classwork: Implement insert().

If the order of unlocks is reversed?

Classwork: Check if two concurrent inserts work.



Classwork: Now allow remove().


```

void insert(Node &prev, Node &naya) {
    prev.lock();
    naya.next = prev.next;
    prev.next = naya;
    prev.unlock();
}

void remove(Node &prev, Node &tbr) {
    tbr.lock();
    prev.lock();
    prev.next = tbr.next;
    // process tbr.
    tbr.unlock();
    prev.unlock();
}

```

If the order of locks is reversed?



```

void insert(Node &prev, Node &naya) {
    prev.lock();
    naya.next = prev.next;
    prev.next = naya;
    prev.unlock();
}

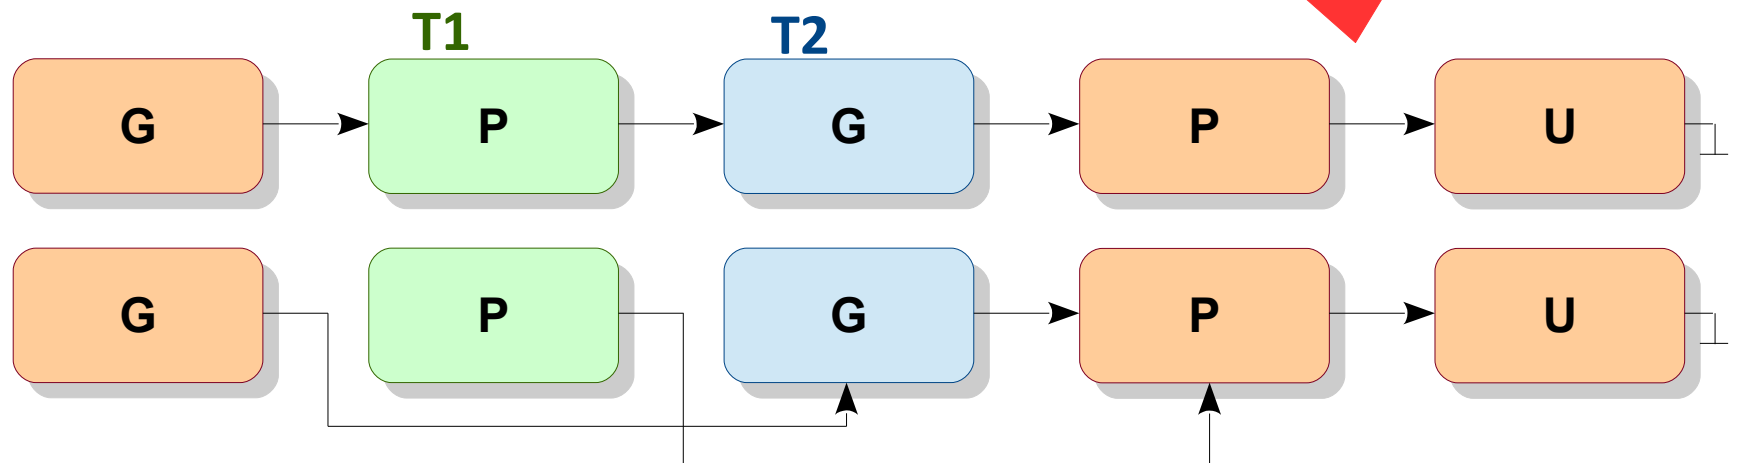
void remove(Node &prev, Node &tbr) {
    prev.lock();
    tbr.lock();
    prev.next = tbr.next;
    // process tbr.
    tbr.unlock();
    prev.unlock();
}

```

Isn't a similar issue possible with our current implementation?

Where is the problem?

Expected GPU,
received GGPU.




```

void remove(Node &prev) {
    prev.lock();
    tbr = prev.next;
    tbr.lock();
    prev.next = tbr.next;
    // process tbr.
    tbr.unlock();
    prev.unlock();
}

```

This would solve the problem of **tbr**, but what about **remove(P)** and **remove(P)** executing concurrently?

This requires us to check for a node's validity.



```
int remove(Node &prev) {
```

```
    if (prev.valid == 0) return -1;  
    prev.lock();
```

```
    tbr = prev.next;
```

```
    tbr.lock();
```

```
    prev.next = tbr.next;
```

```
    // process tbr.
```

```
    tbr.valid = 0;
```


```
    tbr.unlock();
```

```
    prev.unlock();
```

```
}
```

Checking for validity and locking needs to be **atomic**!

- Memory is not reclaimed!
- Only insert and remove!
- No traversal yet!
- Direct pointer is given!
- Still so many complications!

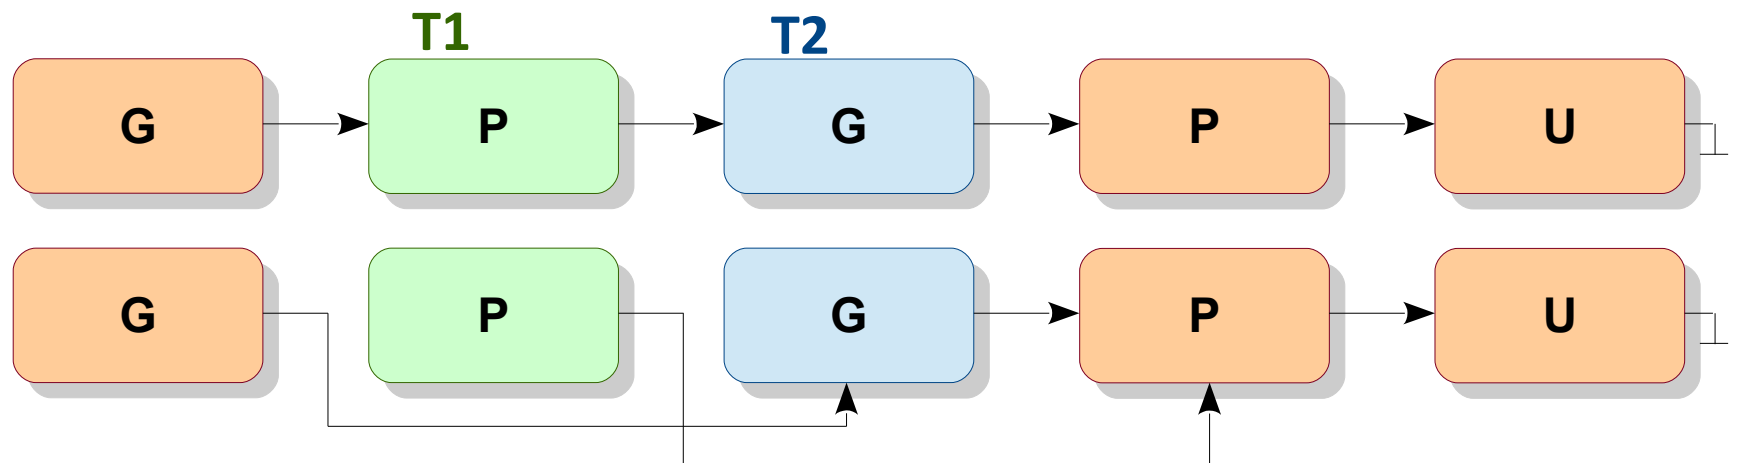


```

int remove(Node &prev) {
    prev.lock();
    if (prev.valid == 0) return -1;
    tbr = prev.next;
    tbr.lock();
    if (tbr.valid == 0) return -2;
    prev.next = tbr.next;
    // process tbr.
    tbr.valid = 0;
    tbr.unlock();
    prev.unlock();
}

```

Does not unlock on error.




```

int remove(Node &prev) {
    prev.lock();
    if (prev.valid) {
        tbr = prev.next;
        tbr.lock();
        if (tbr.valid) {
            prev.next = tbr.next;
            // process tbr.
            tbr.valid = 0;
        }
        tbr.unlock();
    }
    prev.unlock();
}

```

Homework: Find out issues with this code.




Concurrent Linked List

Solution 3: Use atomics to insert.

- Possible only in a few cases
- Difficult to implement with multiple inserts
- Difficult to prove correctness

Classwork: Implement insert().




```

void insert(Node &prev, Node &naya) {
    do {
        naya.next = prev.next;
        old = atomicCAS(&prev.next, naya.next, &naya);
    } while (old != naya.next);
}

```

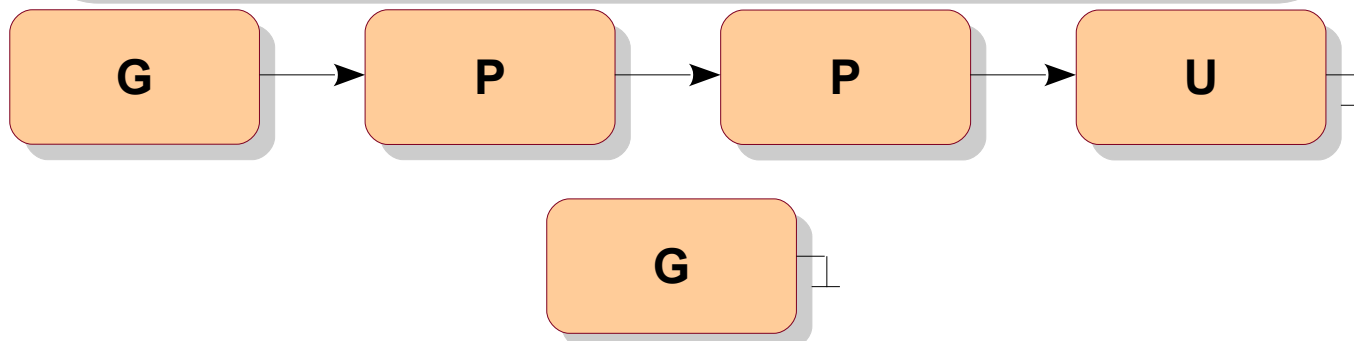
Dare to support remove?

```

void remove(Node &prev) {
    do {
        tbr = prev.next;
        old = atomicCAS(&prev.next, tbr, tbr.next);
    } while (old != tbr);
}

```

Mostly works.
Problem with multiple
concurrent **remove(P)**.



CPU-GPU Synchronization

- While GPU is busy doing work, CPU may perform useful work.
- If CPU-GPU collaborate, they require synchronization.

Classwork: Implement
a functionality to print sequence 0..10.
CPU prints even numbers,
GPU prints odd.

CPU-GPU Synchronization

```
#include <cuda.h>
#include <stdio.h>

__global__ void printk(int *counter) {
    ++*counter;
    printf("\t%d\n", *counter);
}

int main() {
    int hcounter = 0, *counter;

    cudaMalloc(&counter, sizeof(int));

    do {
        printf("%d\n", hcounter);
        cudaMemcpy(counter, &hcounter, sizeof(int), cudaMemcpyHostToDevice);
        printk <<<1, 1>>>(counter);
        cudaMemcpy(&hcounter, counter, sizeof(int), cudaMemcpyDeviceToHost);
    } while (++hcounter < 10);

    return 0;
}
```

Pinned Memory

- Typically, memories are pageable (swappable).
- CUDA allows to make host memory pinned.
- CUDA allows direct access to pinned host memory from device.
- `cudaHostAlloc(&pointer, size, 0);`

Classwork: Implement
the same functionality to print sequence 0..10.
CPU prints even numbers,
GPU prints odd.

Pinned Memory

```
#include <cuda.h>
#include <stdio.h>

__global__ void printk(int *counter) {
    ++*counter;
    printf("\t%d\n", *counter);
}

int main() {
    int *counter;

    cudaHostAlloc(&counter, sizeof(int), 0);

    do {
        printf("%d\n", *counter);
        printk <<<1, 1>>>(counter);
        cudaDeviceSynchronize();
        ++*counter;
    } while (*counter < 10);

    cudaFreeHost(counter);
    return 0;
}
```

No cudaMemcpy!

Classwork: Can we avoid repeated kernel calls?

Persistent Kernels

```
__global__ void printk(int *counter) {  
    do {  
        while (*counter % 2) ;  
        ++*counter;  
        printf("\t%d\n", *counter);  
    } while (*counter < 10);  
}  
int main() {  
    int *counter;  
  
    cudaHostAlloc(&counter, sizeof(int), 0);  
    printk <<<1, 1>>>(counter);  
  
    do {  
        printf("%d\n", *counter);  
        while (*counter % 2 == 0) ;  
        ++*counter;  
    } while (*counter < 10);  
  
    cudaFreeHost(counter);  
    return 0;  
}
```


Extra

Barrier-based Synchronization

→ Disjoint accesses

- Overlapping accesses
- Benign overlaps

Consider threads pushing elements into a worklist



atomic per element



$O(e)$ atomics

atomic per thread




$O(t)$ atomics

prefix-sum




$O(\log t)$ barriers



Barrier-based Synchronization

- Disjoint accesses
- **Overlapping accesses**
- Benign overlaps

Consider threads trying to own a set of elements



atomic per element



e.g., for owning cavities in
Delaunay mesh refinement

non-atomic mark



prioritized mark



check



*Race
and
resolve*

AND

e.g., for inserting unique
elements into a worklist

non-atomic mark



check



*Race
and
resolve*


OR

Barrier-based Synchronization

- Disjoint accesses
 - Overlapping accesses
- **Benign overlaps**

e.g., level-by-level
breadth-first search

Consider threads updating shared
variables to the same value



with atomics



without atomics




Exploiting Algebraic Properties


→ Monotonicity

- Idempotency
- Associativity


Consider threads updating distances in shortest paths computation



Atomic-free update



Lost-update problem




Correction by topology-driven processing, exploiting monotonicity

Exploiting Algebraic Properties

- Monotonicity
- ➔ **Idempotency**
- Associativity


Consider threads updating distances in shortest paths computation



Update by multiple threads



Multiple instances of a node in the worklist




Same node processed by multiple threads

Exploiting Algebraic Properties

- Monotonicity
- Idempotency
- **Associativity**


Consider threads pushing
information to a node



Associativity helps push
information using prefix-sum

Scatter-Gather

Consider threads pushing elements into a worklist



atomic per element  $O(e)$ atomics

atomic per thread  $O(t)$ atomics

prefix-sum  $O(\log t)$ barriers

scatter 

gather 